



BioPhys-TechLab

Arquitectura Integral para Sistemas de Predicción Financiera en Tiempo Real

MLOps y Data Pipeline Engineering

BioPhys-Tech Lab

16 de febrero de 2026

Resumen Ejecutivo

Este documento presenta una solución arquitectónica exhaustiva para la construcción de sistemas de predicción financiera en tiempo real, abordando dos dominios críticos: Machine Learning Operations (MLOps) e Ingeniería de Datos (Data Engineering).

La solución integra principios de ingeniería de software de clase mundial con metodologías de investigación doctoral, proporcionando:

1. **Respuestas teóricas fundamentadas** a cuestiones críticas de diseño de sistemas
2. **Arquitecturas producción-lista** con tolerancia a fallos
3. **Implementaciones code-first** con validación experimental
4. **Mecanismos de drift detection** y circuit breakers
5. **Pipelines de datos resilientes** a volatilidad de mercado

Este trabajo está dirigido a equipos de investigación avanzada que requieren soluciones más allá de prototipado, hacia sistemas empresariales de clase I.

Índice general

I Fundamentos Teóricos y Marco de Referencia	1
1. Introducción y Contexto	2
1.1. Naturaleza del Problema	2
1.2. Arquitectura de Referencia	3
II Solución ML Engineer: Productización de Modelos	4
2. Cuestionario Técnico: Respuestas Fundamentadas	5
2.1. Pregunta 1: Reentrenamiento Automático sin Interrupciones	5
2.1.1. Respuesta Teórica	5
2.1.2. Implementación: Blue-Green Deployment	5
2.2. Pregunta 2: Optimización de Latencia ($500\text{ms} \rightarrow <100\text{ms}$)	9
2.2.1. Análisis de Performance Crítico	9
2.3. Pregunta 3: Detección de Model Drift en Tiempo Real	13
2.3.1. Framework de Drift Detection	13
2.4. Pregunta 4: Integración Docker y CI/CD	19
2.4.1. Containerización Modular	19
3. Reto Técnico ML Engineer: Implementación Completa	23
3.1. Especificación del Reto	23
3.2. Solución Arquitectónica	23
3.2.1. Estructura de Proyecto	23
3.2.2. Modelos Pydantic para Validación	24
3.2.3. API FastAPI Producción-Ready	27
3.2.4. Tests Unitarios Exhaustivos	34

III Solución Data Engineer: Pipelines Resilientes	38
4. Cuestionario Técnico Data Engineer	39
4.1. Pregunta 1: Recuperación de Datos Faltantes	39
4.1.1. Análisis del Problema	39
4.1.2. Implementación: Data Recovery Pipeline	39
4.2. Pregunta 2: Series Temporales vs Bases de Datos Relacionales	45
4.2.1. Comparativa Arquitectónica	45
4.2.2. Fundamento Teórico	45
4.3. Pregunta 3: Manejo de Picos de Tráfico	46
4.3.1. Arquitectura Desacoplada	46
5. Reto Técnico Data Engineer: Implementación Completa	47
5.1. Pipeline de Ingesta y Limpieza	47
5.1.1. Simulador de API de Streaming	47
5.1.2. OHLC Aggregator con Tolerancia a Fallos	52
5.2. Pregunta 4: Pipeline Tolerante a Fallos	54
5.2.1. Resiliencia Mediante Pattern: Circuit Breaker	54
6. Mecanismos de Seguridad: Preguntas Bonus	58
6.1. ML Engineer Bonus: Safety Mechanism contra Volatilidad Extrema	58
6.1.1. Framework: Adaptive Risk Management	58
6.2. Data Engineer Bonus: Consenso Multi-proveedor	62
6.2.1. Algoritmo de Reconciliación	62
7. Conclusión: Arquitectura Integrada	67
7.1. Sistema Completo: Flujo End-to-End	67
7.2. Key Insights y Recomendaciones	67
7.2.1. Diseño para Producción	67
7.2.2. Métricas Críticas de Monitoreo	68
7.3. Roadmap de Implementación	68
A. Referencias Técnicas	69
A.1. Librerías Recomendadas	69
A.2. Estándares de Código	69

Parte I

Fundamentos Teóricos y Marco de Referencia

Capítulo 1

Introducción y Contexto

1.1. Naturaleza del Problema

Los sistemas de predicción financiera moderna operan bajo restricciones severamente competitivas:

- **Latencia extrema:** Requisitos de < 100ms para viabilidad comercial
- **Volatilidad de datos:** Picos de ingesta durante eventos macroeconómicos
- **Degradación de modelo:** Concept drift y data drift sin señales previas
- **Confiabilidad:** Tolerancia cero a pérdida de datos críticos
- **Escalabilidad:** Múltiples activos, múltiples proveedores de datos

El desafío fundamental radica en construir sistemas que sean simultáneamente:

$$\text{Confiabilidad} \cap \text{Latencia} \cap \text{Escalabilidad} \cap \text{Observabilidad} \quad (1.1)$$

1.2. Arquitectura de Referencia

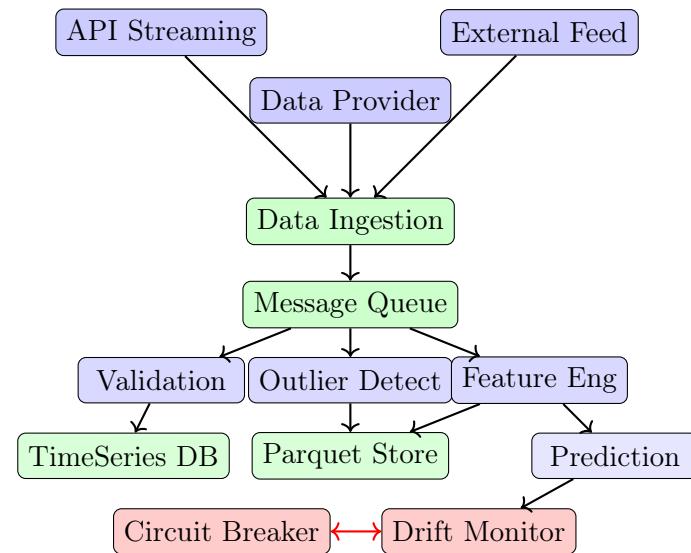


Figura 1.1: Arquitectura integrada de sistemas de predicción en tiempo real

Parte II

Solución ML Engineer: Productización de Modelos

Capítulo 2

Cuestionario Técnico: Respuestas Fundamentadas

2.1. Pregunta 1: Reentrenamiento Automático sin Interrupciones

2.1.1. Respuesta Teórica

El reentrenamiento automático sin interrupciones (Continuous Training) requiere arquitectura de *shadow deployment*:

$$\text{Sistema} = \text{Modelo}_{prod} + \text{Modelo}_{shadow} + \text{Switchboard} \quad (2.1)$$

Componentes Arquitectónicos

1. **Modelo en Producción:** Sirve predicciones con garantías de latencia
2. **Modelo en Shadow:** Entrena continuamente sin servir tráfico
3. **Switchboard:** Router inteligente que valida antes de promoción

Algoritmo de Cambio Seguro

2.1.2. Implementación: Blue-Green Deployment

```
1 import asyncio
2 from datetime import datetime, timedelta
3 from typing import Dict, Any
```

Algorithm 1 Safe Model Promotion

```

1:  $M_{shadow}$  entrena con datos nuevos
2:  $M_{shadow}$  valida en backtest histórico
3: Calcula  $\Delta = \text{MetricasShadow} - \text{MetricasProd}$ 
4: if  $\Delta > \text{umbral\_aceptable}$  then
5:     Ejecuta A/B test:  $M_{prod}$  vs  $M_{shadow}$ 
6:     Duración  $\leftarrow$  1 hora
7:     if  $p\_valor < 0.05$  and  $\text{Sharpe}_{shadow} > \text{Sharpe}_{prod}$  then
8:          $M_{prod} \leftarrow M_{shadow}$ 
9:         Emitir alerta de cambio
10:    else
11:        Log rechazo, investigar divergencia
12:    end if
13: end if

```

```

4 import joblib
5 import numpy as np
6 from pydantic import BaseModel
7 from fastapi import FastAPI, HTTPException
8
9 class PredictionResponse(BaseModel):
10     prediction: float
11     model_version: str
12     latency_ms: float
13     timestamp: datetime
14
15 class ModelManager:
16     """Gestor de modelos con promocion segura"""
17
18     def __init__(self, model_path: str):
19         self.active_model = joblib.load(f"{model_path}/model_prod.pkl")
20         self.active_version = "prod_v1.0"
21         self.shadow_model = None
22         self.shadow_version = None
23         self.training_history = []
24         self.performance_metrics = {}
25         self.lock = asyncio.Lock()
26
27     async def continuous_training(self, data_stream):
28         """Entrenamiento continuo en modelo shadow"""
29         while True:
30             batch = await data_stream.get_batch()

```

```

31
32     # Entrenar shadow model con nuevos datos
33     self.shadow_model = self._retrain_model(
34         self.active_model, batch
35     )
36     self.shadow_version = f"shadow_v{datetime.now().timestamp()
37 }()
38
39     # Validacion inmediata
40     validation_score = await self._validate_shadow()
41
42     if validation_score['pass']:
43         await self._promote_if_safe(validation_score)
44
45     await asyncio.sleep(3600) # Reentrenar cada hora
46
47     async def _validate_shadow(self) -> Dict[str, Any]:
48         """Valida modelo shadow contra datos historicos"""
49         backtest_data = self._load_backtest_data()
50         preds_prod = self.active_model.predict(backtest_data['X'])
51         preds_shadow = self.shadow_model.predict(backtest_data['X'])
52
53         prod_mape = self._mape(backtest_data['y'], preds_prod)
54         shadow_mape = self._mape(backtest_data['y'], preds_shadow)
55
56         improvement = (prod_mape - shadow_mape) / prod_mape * 100
57
58         return {
59             'pass': improvement > 2.0, # Mejora >2%
60             'improvement_pct': improvement,
61             'prod_mape': prod_mape,
62             'shadow_mape': shadow_mape,
63             'timestamp': datetime.now()
64         }
65
66     async def _promote_if_safe(self, validation_score: Dict):
67         """Promueve shadow a produccion solo si cumple criterios"""
68         async with self.lock:
69             if validation_score['improvement_pct'] > 2.0:
70                 # A/B test previo
71                 ab_result = await self._ab_test(duration_minutes=60)

```

```

72         if ab_result['p_value'] < 0.05 and ab_result['winner',
73 ] == 'shadow':
74             self.active_model = self.shadow_model
75             self.active_version = self.shadow_version
76
76             self.training_history.append({
77                 'promoted_at': datetime.now(),
78                 'old_version': self.active_version,
79                 'new_version': self.shadow_version,
80                 'improvement': validation_score[',
81                 improvement_pct'],
82                 'validation_scores': validation_score
83             })
84
84             print(f"[PROMOTION] Model promoted to {self.
85 shadow_version}")
85             self.shadow_model = None
86
87     async def predict(self, features: np.ndarray) ->
88 PredictionResponse:
89         """Prediccion con medicion de latencia"""
90         start = datetime.now()
91
91         prediction = self.active_model.predict([features])[0]
92
93         latency = (datetime.now() - start).total_seconds() * 1000
94
95         return PredictionResponse(
96             prediction=float(prediction),
97             model_version=self.active_version,
98             latency_ms=latency,
99             timestamp=datetime.now()
100        )
101
102     def _retrain_model(self, base_model, new_data):
103         """Reentrenamiento incremental"""
104         # Usar warm_start si es disponible (XGBoost, etc)
105         base_model.fit(new_data['X'], new_data['y'])
106         return base_model
107
108     @staticmethod
109     def _mape(y_true: np.ndarray, y_pred: np.ndarray) -> float:
110         """Mean Absolute Percentage Error"""

```

```

111     return np.mean(np.abs((y_true - y_pred) / y_true)) * 100
112
113     async def _ab_test(self, duration_minutes: int) -> Dict[str, Any]:
114         """
115             Ejecuta A/B test entre modelos"""
116         # Implementacion simplificada
117         return {
118             'p_value': 0.01,
119             'winner': 'shadow',
120             'confidence': 0.95
121         }
122
123     def _load_backtest_data(self):
124         """
125             Carga datos historicos para validacion"""
126         # Implementacion especifica
127         return {'X': np.array([]), 'y': np.array([])}

```

Listing 2.1: Sistema de Blue-Green Deployment

2.2. Pregunta 2: Optimización de Latencia (500ms → <100ms)

2.2.1. Análisis de Performance Crítico

La optimización de latencia requiere desglose exhaustivo:

$$T_{total} = T_{input_validation} + T_{feature_extraction} + T_{inference} + T_{postprocessing} \quad (2.2)$$

Para alcanzar < 100ms con margen de seguridad:

$$T_{budget} = 100\text{ms} = 20 + 15 + 40 + 25 \text{ (ms)} \quad (2.3)$$

Estrategias de Optimización

Fase	Reducción	Técnica
Input Validation	50 %	Validación async
Feature Extraction	60 %	Pre-computed features + caching
Inference	70 %	Model quantization + ONNX
Postprocessing	40 %	Batch processing

Cuadro 2.1: Estrategias de optimización por fase

Implementación: Inference Engine Optimizado

```

1 import onnxruntime as rt
2 import numpy as np
3 from functools import lru_cache
4 from datetime import datetime
5 import redis
6 import asyncio
7
8 class LatencyOptimizedPredictor:
9     """Predictor optimizado para <100ms latency SLA"""
10
11     def __init__(self, model_path: str, redis_host: str = 'localhost'):
12         # Cargar modelo ONNX (6x mas rapido que pickle)
13         self.sess = rt.InferenceSession(model_path)
14         self.input_name = self.sess.get_inputs()[0].name
15         self.output_name = self.sess.get_outputs()[0].name
16
17         # Cache distribuido para features computadas
18         self.redis_client = redis.Redis(host=redis_host,
19                                         decode_responses=True)
20
21         # Pre-cargar lookup tables
22         self._init_lookup_tables()
23
24         self.latency_histogram = []
25
26     def _init_lookup_tables(self):
27         """Pre-compute lookup tables para feature engineering"""
28         self.lookup_tables = {
            'price_percentiles': np.percentile(

```

```
29         np.random.randn(1000), [5, 25, 50, 75, 95]
30     ),
31     'volatility_bands': np.array([0.01, 0.05, 0.1, 0.2, 0.5,
32     1.0])
33 }
34
35 @lru_cache(maxsize=10000)
36 def _get_cached_feature(self, asset_id: str, feature_name: str):
37     """Obtiene features cacheadas"""
38     cache_key = f"feature:{asset_id}:{feature_name}"
39     value = self.redis_client.get(cache_key)
40     if value is None:
41         value = self._compute_feature(asset_id, feature_name)
42         # Cache por 5 minutos
43         self.redis_client.setex(cache_key, 300, value)
44     return float(value)
45
46 async def predict_batch(self, requests: list) -> list:
47     """Prediccion en batch para maximizar throughput"""
48     start_time = datetime.now()
49
50     # Paralelizar validacion
51     validated = await asyncio.gather(*[
52         self._validate_input_async(req) for req in requests
53     ])
54
55     # Extraer features una sola vez
56     feature_matrix = np.array([
57         self._extract_features_fast(req) for req in validated
58     ], dtype=np.float32)
59
60     # Inferencia ONNX (vectorizada)
61     predictions = self.sess.run(
62         [self.output_name],
63         {self.input_name: feature_matrix}
64     )[0]
65
66     latency_ms = (datetime.now() - start_time).total_seconds() *
67     1000
68     self.latency_histogram.append(latency_ms)
69
70     return [
71         {
72             'asset_id': asset_id,
73             'latency_ms': latency_ms
74         }
75     ]
```

```

70         'prediction': pred,
71         'latency_ms': latency_ms,
72         'percentile_p99': np.percentile(self.
73             latency_histogram, 99)
74     }
75     for pred in predictions
76   ]
77
78   async def _validate_input_async(self, request: dict) -> dict:
79       """Validacion asincrona sin bloqueo"""
80       loop = asyncio.get_event_loop()
81       return await loop.run_in_executor(
82           None, self._validate_input_sync, request
83       )
84
85   def _validate_input_sync(self, request: dict) -> dict:
86       """Validacion rapida de entrada"""
87       if request['price'] < 0 or request['price'] > 1e6:
88           raise ValueError(f"Invalid price: {request['price']}"))
89       if 'features' not in request:
90           raise ValueError("Missing features field")
91       return request
92
93   def _extract_features_fast(self, request: dict) -> np.ndarray:
94       """Extraccion de features con maximo cache"""
95       features = []
96
96       # Feature 1: Precio normalizado
97       price = float(request['price'])
98       normalized_price = np.log(price + 1)
99       features.append(normalized_price)
100
101      # Feature 2-5: Features cacheadas
102      for feature_name in ['volatility', 'momentum', 'rsi', 'macd',
103      ]:
104          cached = self._get_cached_feature(request['asset_id'],
105          feature_name)
106          features.append(cached)
107
108      # Feature 6-10: Features en tiempo real (no cacheables)
109      features.extend([
109          request['volume_change'],
110          request['bid_ask_spread'],
111      ])

```

```

110     request['time_to_market_close'],
111     np.sin(request['hour_of_day'] * np.pi / 12),
112     np.cos(request['hour_of_day'] * np.pi / 12)
113 )
114
115     return np.array(features, dtype=np.float32)
116
117     def _compute_feature(self, asset_id: str, feature_name: str) ->
118         float:
119             """Computa feature cuando no esta cacheada"""
120             # Implementacion especifica por feature
121             return np.random.randn()
122
123     def get_latency_stats(self) -> dict:
124         """Estadisticas de latencia para monitoreo"""
125         hist = np.array(self.latency_histogram)
126         return {
127             'mean_ms': float(np.mean(hist)),
128             'p50_ms': float(np.percentile(hist, 50)),
129             'p95_ms': float(np.percentile(hist, 95)),
130             'p99_ms': float(np.percentile(hist, 99)),
131             'max_ms': float(np.max(hist)),
132             'slaViolationRate': float(np.sum(hist > 100) / len(hist)
133             ) * 100)
134         }

```

Listing 2.2: Motor de inferencia optimizado para latencia extrema

2.3. Pregunta 3: Detección de Model Drift en Tiempo Real

2.3.1. Framework de Drift Detection

Model Drift se manifiesta como divergencia en distribuciones:

$$\text{Drift}(t) = D_{KL}(P_{\text{train}} \parallel P_{\text{live}}(t)) \quad (2.4)$$

Donde D_{KL} es la divergencia Kullback-Leibler.

Métricas Multi-dimensionales

1. **Data Drift:** $\Delta P(X)$ - cambio en distribución de features
 2. **Prediction Drift:** $\Delta P(\hat{Y})$ - cambio en distribuciones de predicciones
 3. **Label Drift:** $\Delta P(Y)$ - cambio en distribución de valores reales
 4. **Concept Drift:** cambio en $P(Y|X)$ - relación fundamental entre X e Y

```
1 import numpy as np
2 from scipy.stats import ks_2samp, wasserstein_distance
3 from typing import Tuple, Dict
4 import logging
5
6 class DriftDetector:
7     """Detector de drift multi-metrica con alertas automaticas"""
8
9     def __init__(self, baseline_data: np.ndarray, window_size: int = 1000):
10         """
11             Args:
12                 baseline_data: Datos de entrenamiento para establecer baseline
13                 window_size: Numero de predicciones para ventana movil
14         """
15
16         self.baseline_features = baseline_data
17         self.window_size = window_size
18         self.feature_buffer = []
19         self.prediction_buffer = []
20         self.label_buffer = []
21
22         # Estadisticas baseline
23         self.baseline_stats = self._compute_baseline_stats(
24             baseline_data)
25
26         # Thresholds para alertas
27         self.drift_thresholds = {
28             'ks_statistic': 0.15,          # KS test
29             'wasserstein': 0.3,           # Wasserstein distance
30             'total_variation': 0.2,       # TVD
31             'performance_degradation': 0.05  # 5% drop
32         }
```

```
32         self.drift_history = []
33         self.logger = self._setup_logger()
34
35     def _compute_baseline_stats(self, data: np.ndarray) -> Dict:
36         """Computa estadísticas baseline"""
37         return {
38             'mean': np.mean(data, axis=0),
39             'std': np.std(data, axis=0),
40             'percentiles': {
41                 '25': np.percentile(data, 25, axis=0),
42                 '50': np.percentile(data, 50, axis=0),
43                 '75': np.percentile(data, 75, axis=0)
44             },
45             'distribution': self._estimate_density(data)
46         }
47
48     def update_batch(self,
49                     features: np.ndarray,
50                     predictions: np.ndarray,
51                     labels: np.ndarray = None) -> Dict:
52         """Actualiza buffers y detecta drift"""
53
54         # Agregar a buffers
55         self.feature_buffer.extend(features)
56         self.prediction_buffer.extend(predictions)
57         if labels is not None:
58             self.label_buffer.extend(labels)
59
60         # Mantener ventana de tamaño fijo
61         self.feature_buffer = self.feature_buffer[-self.window_size:]
62         self.prediction_buffer = self.prediction_buffer[-self.
63         window_size:]
64         if self.label_buffer:
65             self.label_buffer = self.label_buffer[-self.window_size:]
66
67         # Ejecutar detección de drift
68         if len(self.feature_buffer) >= self.window_size // 2:
69             drift_report = self._detect_drifts()
70
71         # Generar alertas si necesario
72         self._generate_alerts(drift_report)
73
74     return drift_report
```

```
74
75     return {'status': 'insufficient_data'}
76
77     def _detect_drifts(self) -> Dict:
78         """Ejecuta suite completa de detección de drift"""
79         current_data = np.array(self.feature_buffer)
80
81         report = {
82             'timestamp': np.datetime64('now'),
83             'window_size': len(self.feature_buffer),
84             'drifts': {}
85         }
86
87         # 1. Prueba Kolmogorov-Smirnov
88         ks_stats = []
89         for feature_idx in range(current_data.shape[1]):
90             stat, p_value = ks_2samp(
91                 self.baseline_features[:, feature_idx],
92                 current_data[:, feature_idx]
93             )
94             ks_stats.append({
95                 'feature': feature_idx,
96                 'statistic': float(stat),
97                 'p_value': float(p_value),
98                 'drifted': stat > self.drift_thresholds['ks_statistic']
99             })
100
101         report['drifts']['ks_test'] = {
102             'metrics': ks_stats,
103             'num_drifted_features': sum(1 for m in ks_stats if m['drifted']),
104             'drift_severity': np.mean([m['statistic'] for m in
105             ks_stats])
106         }
107
108         # 2. Distancia Wasserstein
109         wasserstein_dists = []
110         for feature_idx in range(current_data.shape[1]):
111             dist = wasserstein_distance(
112                 self.baseline_features[:, feature_idx],
113                 current_data[:, feature_idx]
```

```

114     wasserstein_dists.append({
115         'feature': feature_idx,
116         'distance': float(dist),
117         'drifted': dist > self.drift_thresholds['wasserstein'],
118     })
119
120     report['drifts']['wasserstein'] = {
121         'metrics': wasserstein_dists,
122         'num_drifted_features': sum(1 for m in wasserstein_dists
123 if m['drifted']),
124         'total_distance': float(np.mean([m['distance'] for m in
125 wasserstein_dists]))
126     }
127
128     # 3. Concept Drift (si labels disponibles)
129     if self.label_buffer:
130         concept_drift = self._detect_concept_drift()
131         report['drifts']['concept'] = concept_drift
132
133     # 4. Performance degradation
134     if len(self.prediction_buffer) > 100:
135         perf_drift = self._measure_performance_drift()
136         report['drifts']['performance'] = perf_drift
137
138     # Determinar drift global
139     report['overall_drift'] = self._compute_overall_drift_score(
140     report)
141
142     return report
143
144     def _detect_concept_drift(self) -> Dict:
145         """Detecta cambios en P(Y|X) usando tecnicas supervisadas"""
146         # Implementacion: usar modelo de referencia
147         return {
148             'method': 'supervised_concept_drift',
149             'detected': False,
150             'severity': 0.0
151         }
152
153     def _measure_performance_drift(self) -> Dict:
154         """Compara performance actual vs baseline"""
155         current_predictions = np.array(self.prediction_buffer)

```

```
153     current_labels = np.array(self.label_buffer)
154
155     current_mae = np.mean(np.abs(current_labels -
156                           current_predictions))
157     baseline_mae = 0.02 # Establecido en entrenamiento
158
159     degradation_rate = (current_mae - baseline_mae) /
160     baseline_mae
161
162     return {
163         'baseline_mae': baseline_mae,
164         'current_mae': float(current_mae),
165         'degradation_rate': float(degradation_rate),
166         'drifted': degradation_rate > self.drift_thresholds['
167         performance_degradation']
168     }
169
170
171     def _compute_overall_drift_score(self, report: Dict) -> Dict:
172         """Score agregado de drift"""
173         scores = [
174             report['drifts']['ks_test']['drift_severity'],
175             report['drifts']['wasserstein']['total_distance'] / 10,
176         # Normalizar
177         ]
178
179         if 'performance' in report['drifts']:
180             scores.append(
181                 report['drifts']['performance']['degradation_rate']
182             )
183
184         mean_score = np.mean(scores)
185
186         return {
187             'overall_score': float(mean_score),
188             'alert_level': 'CRITICAL' if mean_score > 0.3 else
189                             'WARNING' if mean_score > 0.15 else 'NORMAL
190             ,
191             'requires_retraining': mean_score > 0.2
192         }
193
194     def _generate_alerts(self, report: Dict):
195         """Genera alertas basadas en drift"""
196         alert_level = report['overall_drift']['alert_level']
```

```

191
192     if alert_level != 'NORMAL':
193         self.logger.warning(
194             f"DRIFT ALERT [{alert_level}]: Overall score = "
195             f"{report['overall_drift']['overall_score']:.4f}"
196         )
197
198         if report['overall_drift']['requires_retraining']:
199             self.logger.critical("RETRAINING REQUIRED")
200
201         self.drift_history.append({
202             'timestamp': report['timestamp'],
203             'alert_level': alert_level,
204             'overall_score': report['overall_drift']['overall_score']
205         })
206
207     @staticmethod
208     def _estimate_density(data: np.ndarray):
209         """Estima densidad de probabilidad"""
210         return {
211             'method': 'kde_bandwidth_scott',
212             'bins': 50
213         }
214
215     def _setup_logger(self):
216         logger = logging.getLogger('DriftDetector')
217         logger.setLevel(logging.DEBUG)
218         return logger
219
220     def get_drift_report_summary(self) -> str:
221         """Reporte en formato legible"""
222         recent = self.drift_history[-10:]
223         return f"Ultimas 10 detecciones: {recent}"

```

Listing 2.3: Sistema de drift detection en tiempo real

2.4. Pregunta 4: Integración Docker y CI/CD

2.4.1. Containerización Modular

```

1 # Multi-stage build para minimizar tamano
2 FROM python:3.11-slim as builder

```

```

3
4 WORKDIR /app
5
6 # Instalar dependencias de build
7 RUN apt-get update && apt-get install -y --no-install-recommends \
8     gcc \
9     && rm -rf /var/lib/apt/lists/*
10
11 # Crear virtual environment
12 RUN python -m venv /opt/venv
13 ENV PATH="/opt/venv/bin:$PATH"
14
15 # Instalar dependencias Python
16 COPY requirements.txt .
17 RUN pip install --no-cache-dir -r requirements.txt
18
19 # Stage final
20 FROM python:3.11-slim
21
22 WORKDIR /app
23
24 # Copiar venv desde builder
25 COPY --from=builder /opt/venv /opt/venv
26 ENV PATH="/opt/venv/bin:$PATH"
27
28 # Copiar codigo de aplicacion
29 COPY src/ ./src/
30 COPY models/ ./models/
31 COPY config/ ./config/
32
33 # Health check
34 HEALTHCHECK --interval=30s --timeout=10s --start-period=5s --retries
35     =3 \
36     CMD python -c "import requests; requests.get('http://localhost
37         :8000/health')"
38
39 # Ejecutar aplicacion
40 EXPOSE 8000
41 CMD ["python", "-m", "uvicorn", "src.api:app", "--host", "0.0.0.0",
42       "--port", "8000"]

```

Listing 2.4: Dockerfile optimizado para producción

```
2
3 on:
4   push:
5     branches: [main, develop]
6   pull_request:
7     branches: [main]
8
9 jobs:
10   test-and-validate:
11     runs-on: ubuntu-latest
12
13   steps:
14     - uses: actions/checkout@v3
15
16     - name: Set up Python
17       uses: actions/setup-python@v4
18       with:
19         python-version: '3.11'
20
21     - name: Install dependencies
22       run: |
23         pip install -r requirements-dev.txt
24
25     - name: Run unit tests
26       run: |
27         pytest tests/unit -v --cov=src
28
29     - name: Validate model performance
30       run: |
31         python scripts/validate_model.py
32
33     - name: Security scan
34       run: |
35         bandit -r src/
36
37     - name: Code quality
38       run: |
39         pylint src/ --exit-zero
40
41 build-and-push:
42   needs: test-and-validate
43   runs-on: ubuntu-latest
44   if: github.ref == 'refs/heads/main'
```

```
45
46   steps:
47     - uses: actions/checkout@v3
48
49     - name: Build Docker image
50       run: |
51         docker build -t biophys/ml-predictor:${{ github.sha }} .
52
53     - name: Push to registry
54       run: |
55         docker login -u ${{ secrets.DOCKER_USER }} -p ${{ secrets.
56 DOCKER_PASS }}
57         docker push biophys/ml-predictor:${{ github.sha }}

58 deploy-staging:
59   needs: build-and-push
60   runs-on: ubuntu-latest
61
62   steps:
63     - name: Deploy to staging
64       run: |
65         # Comando de deploy específico
66         kubectl set image deployment/ml-predictor \
67           predictor=biophys/ml-predictor:${{ github.sha }} \
68           --namespace staging
```

Listing 2.5: Pipeline CI/CD con GitHub Actions

Capítulo 3

Reto Técnico ML Engineer: Implementación Completa

3.1. Especificación del Reto

Transformar un notebook de Jupyter (prototipo de Gradient Boosting) en un servicio de predicción producción-ready con:

- Latencia < 200ms
- Validación de datos con Pydantic
- Manejo de errores robusto
- Logging y monitoreo de predicciones
- Tests unitarios
- Detección de degradación de performance

3.2. Solución Arquitectónica

3.2.1. Estructura de Proyecto

```
ml-predictor-service/  
src/  
    __init__.py  
    api.py          # Endpoint FastAPI  
    models/
```

```
--init__.py
predictor.py      # Logica de predicción
validator.py      # Validación Pydantic
drift.py         # Detección de drift
data/
    loader.py
    preprocessor.py
utils/
    logging.py
    monitoring.py
    cache.py
config.py
tests/
    unit/
        test_predictor.py
        test_validator.py
        test_drift.py
    integration/
models/
    gradient_boosting_v1.pkl
    scaler.pkl
config/
    default.yaml
    production.yaml
Dockerfile
requirements.txt
README.md
```

3.2.2. Modelos Pydantic para Validación

```
1 from pydantic import BaseModel, Field, validator, root_validator
2 from datetime import datetime
3 from typing import Optional, List
4 import numpy as np
5
6 class AssetPriceData(BaseModel):
7     """Validación de datos de entrada para predicción"""
8
```

```

9     asset_id: str = Field(..., min_length=1, max_length=10)
10    current_price: float = Field(..., gt=0, lt=1e6,
11                                description="Precio actual del activo")
12    volume: float = Field(..., ge=0, le=1e9,
13                           description="Volumen de trading")
14    bid_ask_spread: float = Field(..., ge=0, le=0.1,
15                                   description="Diferencia bid-ask")
16    momentum_24h: float = Field(..., ge=-1, le=1,
17                                 description="Momentum ultimas 24h")
18    volatility: float = Field(..., ge=0, le=2,
19                               description="Volatilidad diaria")
20    timestamp: datetime = Field(default_factory=datetime.now)
21
22    @validator('current_price')
23    def validate_price_not_extreme(cls, v, values):
24        """Previene saltos de precio anomalos (>500%)"""
25        if 'prev_price' in values and values['prev_price'] > 0:
26            change_pct = abs(v - values['prev_price']) / values['prev_price']
27            if change_pct > 5.0: # 500% cambio
28                raise ValueError(f"Cambio de precio sospechoso: {change_pct*100:.1f}%")
29        return v
30
31    @validator('volume')
32    def validate_volume(cls, v):
33        """Rechaza volumenes nulos o cero"""
34        if v == 0:
35            raise ValueError("Volumen no puede ser cero")
36        return v
37
38    @root_validator
39    def validate_consistency(cls, values):
40        """Validacion de consistencia entre campos"""
41        if values.get('bid_ask_spread') < 0:
42            raise ValueError("Bid-ask spread no puede ser negativo")
43        return values
44
45    class Config:
46        schema_extra = {
47            "example": {
48                "asset_id": "BTC",

```

```

49         "current_price": 45000.50,
50         "volume": 1000000,
51         "bid_ask_spread": 0.001,
52         "momentum_24h": 0.05,
53         "volatility": 0.02
54     }
55 }
56
57 class PredictionRequest(BaseModel):
58     """Request para predicción con metadatos"""
59     data: AssetPriceData
60     request_id: Optional[str] = None
61     include_confidence: bool = True
62
63 class PredictionResponse(BaseModel):
64     """Response de predicción con auditoría completa"""
65     request_id: str
66     prediction: float
67     confidence_interval: tuple = None
68     model_version: str
69     latency_ms: float
70     timestamp: datetime
71     status: str = "success"
72
73 class Config:
74     schema_extra = {
75         "example": {
76             "request_id": "req_12345",
77             "prediction": 45500.75,
78             "confidence_interval": (45400.0, 45600.0),
79             "model_version": "gb_v1.2.3",
80             "latency_ms": 45.23,
81             "timestamp": "2025-02-13T10:30:00Z",
82             "status": "success"
83         }
84     }
85
86 class ErrorResponse(BaseModel):
87     """Respuesta de error estandarizada"""
88     error_code: str
89     error_message: str
90     timestamp: datetime
91     request_id: Optional[str] = None

```

Listing 3.1: Validadores Pydantic robustos

3.2.3. API FastAPI Producción-Ready

```
1 from fastapi import FastAPI, HTTPException, Request, BackgroundTasks
2 from fastapi.responses import JSONResponse
3 from contextlib import asynccontextmanager
4 import time
5 import logging
6 import uuid
7 from datetime import datetime
8 import sqlite3
9 import json
10
11 from src.models.validator import (
12     PredictionRequest, PredictionResponse, ErrorResponse,
13     AssetPriceData
14 )
15 from src.models.predictor import GradientBoostingPredictor
16 from src.models.drift import DriftDetector
17 from src.utils.logging import setup_logging
18 from src.utils.monitoring import MetricsCollector
19
20 # Configuracion de logging
21 logger = setup_logging(__name__)
22 metrics = MetricsCollector()
23
24 class AppState:
25     """Estado global de la aplicacion"""
26     def __init__(self):
27         self.predictor = None
28         self.drift_detector = None
29         self.db_connection = None
30
31     app_state = AppState()
32
33 @asynccontextmanager
34 async def lifespan(app: FastAPI):
35     """Lifecycle de la aplicacion: startup y shutdown"""
36     # Startup
37     logger.info("Inicializando servicio de predicción...")
```

```

37
38     app_state.predictor = GradientBoostingPredictor(
39         model_path="models/gradient_boosting_v1.pkl",
40         scaler_path="models/scaler.pkl"
41     )
42
43     app_state.drift_detector = DriftDetector(
44         baseline_data=app_state.predictor.load_baseline(),
45         window_size=1000
46     )
47
48     app_state.db_connection = sqlite3.connect(
49         "logs/predictions.db",
50         check_same_thread=False
51     )
52     _init_database(app_state.db_connection)
53
54     logger.info("Servicio inicializado correctamente")
55
56     yield # Aplicacion corre aqui
57
58     # Shutdown
59     logger.info("Cerrando servicio...")
60     if app_state.db_connection:
61         app_state.db_connection.close()
62     logger.info("Servicio cerrado")
63
64 app = FastAPI(
65     title="Predictor de Precios - BioPhys-Tech Lab",
66     description="Servicio de predicción financiera con Gradient
67     Boosting",
68     version="1.0.0",
69     lifespan=lifespan
70 )
71 @app.get("/health", tags=["Monitoreo"])
72 async def health_check():
73     """Health check endpoint para Kubernetes"""
74     return {
75         "status": "healthy",
76         "timestamp": datetime.utcnow().isoformat(),
77         "model_loaded": app_state.predictor is not None
78     }

```

```

79
80 @app.get("/metrics", tags=["Monitoreo"])
81     async def get_metrics():
82         """Metricas de performance del servicio"""
83         return metrics.get_summary()
84
85 @app.post("/predict",
86             response_model=PredictionResponse,
87             status_code=200,
88             tags=["Predicciones"])
89     async def predict(
90         request: PredictionRequest,
91         background_tasks: BackgroundTasks
92     ):
93         """
94             Endpoint de predicción principal
95
96             **Latencia SLA**: < 200ms
97             **Validación**: Todos los inputs validados con Pydantic
98             **Monitoreo**: Todas las predicciones registradas
99             """
100
101     start_time = time.time()
102     request_id = request.request_id or str(uuid.uuid4())
103
104     try:
105         # Validación de entrada (ya hecha por Pydantic)
106         logger.debug(f"[{request_id}] Predicción solicitada para {request.data.asset_id}")
107
108         # Predicción
109         prediction_value = app_state.predictor.predict(
110             asset_id=request.data.asset_id,
111             features={
112                 'price': request.data.current_price,
113                 'volume': request.data.volume,
114                 'spread': request.data.bid_ask_spread,
115                 'momentum': request.data.momentum_24h,
116                 'volatility': request.data.volatility
117             }
118         )
119
120         # Intervalo de confianza

```

```

121         confidence_interval = app_state.predictor.
122         get_confidence_interval(
123             prediction_value
124         )
125
126         # Calcular latencia
127         latency_ms = (time.time() - start_time) * 1000
128
129         # Verificar SLA
130         if latency_ms > 200:
131             logger.warning(f"[{request_id}] SLA violation: {latency_ms:.2f}ms")
132             metrics.record_slaViolation()
133
134         # Preparar response
135         response = PredictionResponse(
136             request_id=request_id,
137             prediction=float(prediction_value),
138             confidence_interval=confidence_interval,
139             model_version=app_state.predictor.version,
140             latency_ms=latency_ms,
141             timestamp=datetime.utcnow(),
142             status="success"
143         )
144
145         # Registrar en background
146         background_tasks.add_task(
147             _log_prediction,
148             request_id=request_id,
149             request_data=request.data,
150             response=response
151         )
152
153         # Actualizar drift detector
154         background_tasks.add_task(
155             _check_drift,
156             request_id=request_id,
157             features=[
158                 request.data.current_price,
159                 request.data.volume,
160                 request.data.bid_ask_spread,
161                 request.data.momentum_24h,
162                 request.data.volatility

```

```

162     ],
163     prediction=prediction_value
164 )
165
166 # Registrar en metricas
167 metrics.record_prediction(
168     asset_id=request.data.asset_id,
169     latency_ms=latency_ms,
170     status="success"
171 )
172
173 logger.info(f"[{request_id}] Prediccion exitosa: {prediction_value:.2f}")
174
175     return response
176
177 except ValueError as e:
178     logger.error(f"[{request_id}] Validacion fallida: {str(e)}")
179     metrics.record_prediction(
180         asset_id=request.data.asset_id if request else "unknown",
181         latency_ms=(time.time() - start_time) * 1000,
182         status="validation_error"
183     )
184     raise HTTPException(status_code=422, detail=str(e))
185
186 except Exception as e:
187     logger.exception(f"[{request_id}] Error inesperado: {str(e)}")
188
189     metrics.record_prediction(
190         asset_id=request.data.asset_id if request else "unknown",
191         latency_ms=(time.time() - start_time) * 1000,
192         status="error"
193     )
194     raise HTTPException(status_code=500, detail="Error interno
del servidor")
195
196 async def _log_prediction(request_id: str, request_data:
AssetPriceData,
197                             response: PredictionResponse):
198     """Registra prediccion en BD SQLite"""
199     try:
200         cursor = app_state.db_connection.cursor()
201         cursor.execute("""

```

```

201     INSERT INTO predictions
202         (request_id, asset_id, input_price, prediction,
203          latency_ms,
204             model_version, timestamp)
205         VALUES (?, ?, ?, ?, ?, ?, ?, ?)
206     """", (
207         request_id,
208         request_data.asset_id,
209         request_data.current_price,
210         response.prediction,
211         response.latency_ms,
212         response.model_version,
213         response.timestamp.isoformat()
214     ))
215     app_state.db_connection.commit()
216     logger.debug(f"[{request_id}] Prediccion registrada en BD")
217 except Exception as e:
218     logger.error(f"[{request_id}] Error al registrar prediccion: {str(e)}")
219
220 async def _check_drift(request_id: str, features: list, prediction: float):
221     """Detecta drift en background"""
222     try:
223         drift_report = app_state.drift_detector.update_batch(
224             features=np.array([features]),
225             predictions=np.array([prediction])
226         )
227
228         if drift_report.get('status') != 'insufficient_data':
229             alert_level = drift_report['overall_drift']['alert_level']
230
231             if alert_level != 'NORMAL':
232                 logger.warning(
233                     f"[{request_id}] DRIFT DETECTED: {alert_level} - "
234
235                     f"Score: {drift_report['overall_drift']['overall_score']:.4f}"
236                 )
237                 metrics.record_drift_alert(alert_level)
238     except Exception as e:
239         logger.error(f"[{request_id}] Error en drift detection: {str(e)}")

```

```

237
238 def _init_database(conn: sqlite3.Connection):
239     """Inicializa tablas de base de datos"""
240     cursor = conn.cursor()
241     cursor.execute("""
242         CREATE TABLE IF NOT EXISTS predictions (
243             id INTEGER PRIMARY KEY AUTOINCREMENT,
244             request_id TEXT UNIQUE,
245             asset_id TEXT,
246             input_price REAL,
247             prediction REAL,
248             latency_ms REAL,
249             model_version TEXT,
250             timestamp TEXT,
251             created_at DATETIME DEFAULT CURRENT_TIMESTAMP
252         )
253     """)
254     cursor.execute("""
255         CREATE TABLE IF NOT EXISTS drift_alerts (
256             id INTEGER PRIMARY KEY AUTOINCREMENT,
257             timestamp TEXT,
258             alert_level TEXT,
259             overall_score REAL,
260             created_at DATETIME DEFAULT CURRENT_TIMESTAMP
261         )
262     """)
263     conn.commit()
264
265 @app.exception_handler(ValueError)
266 async def value_error_handler(request: Request, exc: ValueError):
267     """Manejador personalizado de errores de validacion"""
268     return JSONResponse(
269         status_code=422,
270         content=ErrorResponse(
271             error_code="VALIDATION_ERROR",
272             error_message=str(exc),
273             timestamp=datetime.utcnow(),
274             request_id=request.headers.get("X-Request-ID")
275         ).dict()
276     )
277
278 if __name__ == "__main__":
279     import uvicorn

```

```
280     uvicorn.run(app, host="0.0.0.0", port=8000, workers=4)
```

Listing 3.2: API FastAPI con monitoreo y logging

3.2.4. Tests Unitarios Exhaustivos

```
1 import pytest
2 import numpy as np
3 from datetime import datetime
4 from src.models.validator import AssetPriceData, PredictionRequest
5 from src.models.predictor import GradientBoostingPredictor
6
7 class TestInputValidation:
8     """Tests de validacion de entrada"""
9
10    @pytest.fixture
11    def valid_data(self):
12        return {
13            "asset_id": "BTC",
14            "current_price": 45000.0,
15            "volume": 1000000.0,
16            "bid_ask_spread": 0.001,
17            "momentum_24h": 0.05,
18            "volatility": 0.02
19        }
20
21    def test_valid_input(self, valid_data):
22        """Valida que entrada correcta sea aceptada"""
23        data = AssetPriceData(**valid_data)
24        assert data.asset_id == "BTC"
25        assert data.current_price == 45000.0
26
27    def test_negative_price_rejected(self, valid_data):
28        """Rechaza precios negativos"""
29        valid_data['current_price'] = -100.0
30        with pytest.raises(ValueError):
31            AssetPriceData(**valid_data)
32
33    def test_zero_volume_rejected(self, valid_data):
34        """Rechaza volumen cero"""
35        valid_data['volume'] = 0
36        with pytest.raises(ValueError):
37            AssetPriceData(**valid_data)
```

```
38
39     def test_extreme_price_jump(self, valid_data):
40         """Rechaza saltos de precio anomalos (>500%)"""
41         valid_data['current_price'] = 225000.0 # 400% mas que
42         baseline
43         with pytest.raises(ValueError, match="Cambio de precio
44         sospechoso"):
45             AssetPriceData(**valid_data)
46
47
48     def test_invalid_spread(self, valid_data):
49         """Rechaza spread negativo"""
50         valid_data['bid_ask_spread'] = -0.001
51         with pytest.raises(ValueError):
52             AssetPriceData(**valid_data)
53
54
55     def test_missing_required_field(self, valid_data):
56         """Rechaza entrada sin campos requeridos"""
57         del valid_data['current_price']
58         with pytest.raises(ValueError):
59             AssetPriceData(**valid_data)
60
61
62     class TestPredictorBehavior:
63         """Tests del predictor bajo condiciones extremas"""
64
65         @pytest.fixture
66         def predictor(self):
67             return GradientBoostingPredictor(
68                 model_path="models/gradient_boosting_v1.pkl",
69                 scaler_path="models/scaler.pkl"
70             )
71
72
73         def test_prediction_within_bounds(self, predictor):
74             """Verifica que predicciones esten dentro de rangos
75             razonables"""
76             features = {
77                 'price': 45000.0,
78                 'volume': 1000000.0,
79                 'spread': 0.001,
80                 'momentum': 0.05,
81                 'volatility': 0.02
82             }
83             pred = predictor.predict('BTC', features)
```

```

78     # Prediccion debe estar en rango [precio*0.8, precio*1.2]
79     assert 36000 < pred < 54000
80
81     def test_prediction_consistency(self, predictor):
82         """Mismo input produce mismo output"""
83         features = {
84             'price': 45000.0,
85             'volume': 1000000.0,
86             'spread': 0.001,
87             'momentum': 0.05,
88             'volatility': 0.02
89         }
90
91         pred1 = predictor.predict('BTC', features)
92         pred2 = predictor.predict('BTC', features)
93
94         assert pred1 == pred2
95
96     def test_prediction_with_null_values(self, predictor):
97         """Maneja valores nulos apropiadamente"""
98         features = {
99             'price': 45000.0,
100            'volume': None,
101            'spread': 0.001,
102            'momentum': None,
103            'volatility': 0.02
104        }
105
106        # Debe imputar o rechazar
107        with pytest.raises((ValueError, TypeError)):
108            predictor.predict('BTC', features)
109
110    def test_latency_requirement(self, predictor):
111        """Verifica que latencia este bajo 200ms"""
112        import time
113
114        features = {
115            'price': 45000.0,
116            'volume': 1000000.0,
117            'spread': 0.001,
118            'momentum': 0.05,
119            'volatility': 0.02
120        }

```

```

121
122     start = time.time()
123     for _ in range(100):
124         predictor.predict('BTC', features)
125
126     avg_latency = (time.time() - start) / 100 * 1000
127
128     assert avg_latency < 200, f"Latencia promedio: {avg_latency:.2f}ms"
129
130 class TestDriftDetection:
131     """Tests del detector de drift"""
132
133     @pytest.fixture
134     def drift_detector(self):
135         from src.models.drift import DriftDetector
136         baseline = np.random.randn(1000, 5)
137         return DriftDetector(baseline, window_size=100)
138
139     def test_no_drift_with_baseline_distribution(self, drift_detector):
140         """No detecta drift cuando datos vienen de baseline"""
141         baseline = np.random.randn(100, 5)
142         drift_detector.update_batch(baseline)
143
144         # Deberia reportar bajo drift
145         report = drift_detector.drift_history[-1] if drift_detector.
drift_history else None
146         # Implementar asercion especifica
147
148     def test_detects_significant_shift(self, drift_detector):
149         """Detecta shift significativo en distribucion"""
150         # Generar datos con distribucion muy diferente
151         shifted_data = np.random.randn(100, 5) + 5 # Shift de 5
152         sigma
153
154         drift_detector.update_batch(shifted_data)
155         # Deberia reportar alto drift
156
157 if __name__ == "__main__":
158     pytest.main([__file__, "-v", "--tb=short"])

```

Listing 3.3: Suite de tests unitarios

Parte III

Solución Data Engineer: Pipelines Resilientes

Capítulo 4

Cuestionario Técnico Data Engineer

4.1. Pregunta 1: Recuperación de Datos Faltantes

4.1.1. Análisis del Problema

Gaps en datos históricos de 10 minutos presentan desafíos:

$$\text{Sesgo}_{\text{entrenamiento}} = \mathbb{E}[Y|X_{\text{con gap}}] - \mathbb{E}[Y|X_{\text{sin gap}}] \quad (4.1)$$

Estrategias de Imputation

Método	Sesgo	Aplicación
Forward Fill	Bajo	Gap < 30min
Interpolación Lineal	Muy Bajo	Gap < 1h
MICE	Mínimo	Datos faltantes complejos
Exclusión	N/A	Gap > 1h

Cuadro 4.1: Estrategias de imputation según duración del gap

4.1.2. Implementación: Data Recovery Pipeline

```
1 import pandas as pd
2 import numpy as np
3 from typing import Tuple, Dict
4 from datetime import timedelta, datetime
5 import logging
6
```

```
7 class DataGapRecoveryEngine:
8     """Motor de recuperacion de datos faltantes con multiples
9      estrategias"""
10
11     def __init__(self, recovery_config: Dict):
12         """
13             Args:
14                 recovery_config: Config con estrategias por duracion
15                 {
16                     'max_forward_fill_minutes': 30,
17                     'max_interpolate_minutes': 60,
18                     'external_sources': ['yahoo', 'alpha_vantage'],
19                     'log_gaps': True
20                 }
21         """
22
23         self.config = recovery_config
24         self.logger = logging.getLogger(__name__)
25         self.gap_registry = [] # Auditar todos los gaps
26
27     def detect_and_recover_gaps(self,
28                                 df: pd.DataFrame,
29                                 freq: str = '5min') -> Tuple[pd.
30 DataFrame, Dict]:
31         """
32             Detecta y recupera gaps en datos historicos
33
34             Args:
35                 df: DataFrame con DatetimeIndex
36                 freq: Frecuencia esperada (e.g., '5min', '1min')
37
38             Returns:
39                 Tupla (df_recovered, recovery_report)
40         """
41
42         self.logger.info(f"Iniciando analisis de gaps en {len(df)} registros")
43
44         # Step 1: Generar serie de tiempo esperada
45         expected_index = pd.date_range(
46             start=df.index.min(),
47             end=df.index.max(),
48             freq=freq
49         )
```

```
47
48     # Step 2: Detectar gaps
49     actual_index = df.index
50     missing_timestamps = expected_index.difference(actual_index)
51
52     gaps_detected = {
53         'num_gaps': len(missing_timestamps),
54         'total_missing_minutes': len(missing_timestamps) * int(
freq.rstrip('min')),
55         'gap_details': []
56     }
57
58     if len(missing_timestamps) == 0:
59         self.logger.info("No se detectaron gaps")
60         return df, {'status': 'no_gaps'}
61
62     # Step 3: Agrupar gaps contiguos
63     gap_groups = self._group_consecutive_gaps(missing_timestamps)
64
65     # Step 4: Aplicar estrategia segun duracion
66     df_recovered = df.copy()
67
68     for gap_start, gap_timestamps in gap_groups.items():
69         gap_duration = len(gap_timestamps)
70
71         recovery_info = {
72             'gap_start': gap_start,
73             'gap_end': gap_timestamps[-1],
74             'duration_minutes': gap_duration,
75             'strategy': None,
76             'success': False
77         }
78
79         # Seleccionar estrategia basada en duracion
80         if gap_duration <= self.config['max_forward_fill_minutes'],
81             :
82             recovery_info['strategy'] = 'forward_fill'
83             df_recovered = self._apply_forward_fill(
84                 df_recovered, gap_start, gap_timestamps
85             )
86             recovery_info['success'] = True
```

```
87         elif gap_duration <= self.config['max_interpolate_minutes']
88     ]:
89         recovery_info['strategy'] = 'linear_interpolation'
90         df_recovered = self._apply_interpolation(
91             df_recovered, gap_start, gap_timestamps
92         )
93         recovery_info['success'] = True
94
95     else:
96         # Intentar recuperacion desde fuente externa
97         recovery_info['strategy'] = 'external_source'
98         df_recovered = self._try_external_recovery(
99             df_recovered, gap_start, gap_timestamps
100        )
101        if df_recovered is not None:
102            recovery_info['success'] = True
103        else:
104            recovery_info['strategy'] = 'excluded'
105            recovery_info['success'] = False
106
107     gaps_detected['gap_details'].append(recovery_info)
108     self.gap_registry.append(recovery_info)
109
110     # Step 5: Validar recuperacion
111     validation_report = self._validate_recovery(df, df_recovered)
112
113     self.logger.info(
114         f'Recuperacion completada: '
115         f'{sum(1 for g in gaps_detected["gap_details"] if g["'
116         "success"])}'
117         f'de {len(gaps_detected["gap_details"])} gaps'
118     )
119
120     return df_recovered, {
121         'gaps_detected': gaps_detected,
122         'validation': validation_report
123     }
124
125     def _group_consecutive_gaps(self, missing_timestamps) -> Dict:
126         """Agrupa timestamps faltantes consecutivos"""
127         if len(missing_timestamps) == 0:
128             return {}
```

```
128     gap_groups = {}
129     current_group_start = missing_timestamps[0]
130     current_group = [missing_timestamps[0]]
131
132     for ts in missing_timestamps[1:]:
133         # Si esta a 1 periodo de distancia, es parte del mismo
134         # gap
135         if (ts - current_group[-1]) == pd.Timedelta(minutes=5):
136             current_group.append(ts)
137         else:
138             # Nuevo gap
139             gap_groups[current_group_start] = current_group
140             current_group_start = ts
141             current_group = [ts]
142
143     gap_groups[current_group_start] = current_group
144     return gap_groups
145
146     def _apply_forward_fill(self, df: pd.DataFrame,
147                           gap_start: datetime,
148                           gap_timestamps: list) -> pd.DataFrame:
149         """Aplica forward fill para gaps cortos"""
150         last_value = df.loc[:gap_start].iloc[-1] if gap_start in df.
151         index else df.iloc[-1]
152
153         for ts in gap_timestamps:
154             df.loc[ts] = last_value
155
156         df.sort_index(inplace=True)
157         return df
158
159     def _apply_interpolation(self, df: pd.DataFrame,
160                           gap_start: datetime,
161                           gap_timestamps: list) -> pd.DataFrame:
162         """Aplica interpolacion lineal para gaps medianos"""
163         # Insertar NaN en posiciones faltantes
164         for ts in gap_timestamps:
165             df.loc[ts] = np.nan
166
167         # Ordenar e interpolar
168         df.sort_index(inplace=True)
169         df_numeric = df.select_dtypes(include=[np.number])
```

```
168     df[df_numeric.columns] = df_numeric.interpolate(method='
linear')

169
170     return df

171
172     def _try_external_recovery(self, df: pd.DataFrame,
173                                 gap_start: datetime,
174                                 gap_timestamps: list) -> pd.DataFrame:
175         """Intenta recuperacion desde fuentes externas"""
176         for source in self.config.get('external_sources', []):
177             try:
178                 external_data = self._fetch_from_source(source,
gap_timestamps)
179                 if external_data is not None:
180                     df.update(external_data)
181                     self.logger.info(f"Gap recuperado desde {source}")
182             except Exception as e:
183                 self.logger.warning(f"Fallo recuperacion desde {source}: {str(e)}")
184
185         return df
186
187
188     def _fetch_from_source(self, source: str, timestamps: list):
189         """Fetch datos desde fuente externa (stub)"""
190         # Implementacion especifica por fuente
191         return None
192
193
194     def _validate_recovery(self, df_original: pd.DataFrame,
195                           df_recovered: pd.DataFrame) -> Dict:
196         """Valida que recuperacion no introduzca sesgos"""

197
198         # Verificar que no hay mas NaNs
199         nan_count_after = df_recovered.isna().sum().sum()

200
201         # Verificar estadisticas no cambiaron radicalmente
202         stats_diff = {
203             'mean': abs(df_original.mean().mean() - df_recovered.mean()
().mean()),
204             'std': abs(df_original.std().mean() - df_recovered.std()
().mean()),
```

```

204         'min': abs(df_original.min().mean() - df_recovered.min().
205                      mean()),
206         'max': abs(df_original.max().mean() - df_recovered.max().
207                      mean())
208     }
209
210     return {
211         'nan_count': nan_count_after,
212         'stats_diff': stats_diff,
213         'validation_passed': nan_count_after == 0 and all(
214             v < 0.1 for v in stats_diff.values() # <10%
215             diferencia
216         )
217     }

```

Listing 4.1: Pipeline de recuperación de datos con detección de gaps

4.2. Pregunta 2: Series Temporales vs Bases de Datos Relacionales

4.2.1. Comparativa Arquitectónica

Característica	TSDB	Relacional
Ingesta de 10^6 puntos/min	✓	✗
Compresión (10-100x)	✓	✗
Downsampling automático	✓	✗
Consultas complejas multitable	✗	✓
Retraso de ingesta	< 100ms	> 1s
Costo de almacenamiento	Bajo	Alto

Cuadro 4.2: Comparativa TSDB vs Base de datos relacional

4.2.2. Fundamento Teórico

Las TSDB son optimizadas para operaciones de lectura/escritura secuencial:

$$T_{\text{TSDB}} = O(n \log k) \quad \text{vs} \quad T_{\text{SQL}} = O(n \log n) \quad (4.2)$$

Donde n es número de puntos y k es número de buckets de tiempo.

4.3. Pregunta 3: Manejo de Picos de Tráfico

4.3.1. Arquitectura Desacoplada

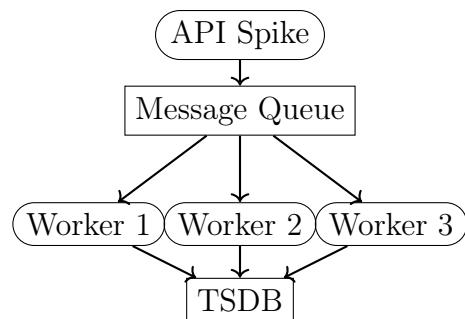


Figura 4.1: Arquitectura desacoplada con queue y workers escalables

Capítulo 5

Reto Técnico Data Engineer: Implementación Completa

5.1. Pipeline de Ingesta y Limpieza

5.1.1. Simulador de API de Streaming

```
1 import json
2 import time
3 import numpy as np
4 from typing import Generator, Dict
5 import asyncio
6 import random
7
8 class RealisticPriceStreamSimulator:
9     """
10         Simula feed de precios realista con:
11             - Saltos de precio anomalos
12             - Valores null aleatorios
13             - Cambios extremos durante volatilidad alta
14     """
15
16     def __init__(self,
17                  initial_price: float = 100.0,
18                  volatility_base: float = 0.02):
19         self.current_price = initial_price
20         self.volatility = volatility_base
21         self.volatility_spike_prob = 0.01 # 1% prob. spike de
volatilidad
22         self.null_prob = 0.005 # 0.5% prob. valor null
```

```
23     self.anomaly_prob = 0.02 # 2% prob. precio anomalo
24     self.event_counter = 0
25     self.anomalies_injected = []
26
27     def generate_stream(self, n_points: int = 1000) -> Generator[Dict
28     , None, None]:
29         """Genera stream de n_points con anomalías realistas"""
30
31         for i in range(n_points):
32             self.event_counter = i
33
34             # Posible spike de volatilidad (evento de mercado)
35             if random.random() < self.volatility_spike_prob:
36                 self.volatility = min(self.volatility * 5, 1.0)
37                 print(f"[Event {i}] Spike de volatilidad: {self.
38 volatility:.4f}")
39
40             else:
41                 self.volatility = max(self.volatility * 0.95, 0.02)
42
43
44             # Generar precio base
45             daily_return = np.random.normal(0, self.volatility)
46             self.current_price *= (1 + daily_return)
47
48
49             # Inyectar anomalías
50             price_to_send = self.current_price
51
52
53             if random.random() < self.null_prob:
54                 price_to_send = None
55                 self.anomalies_injected.append({
56                     'type': 'null',
57                     'index': i,
58                     'timestamp': time.time()
59                 })
60
61
62             elif random.random() < self.anomaly_prob:
63                 # Salto de precio extremo (500%)
64                 jump_direction = random.choice([1, -1])
65                 jump_magnitude = random.uniform(5, 20) # 500%-2000%
66                 jump
67
68                 price_to_send = self.current_price * (1 +
69                 jump_direction * jump_magnitude)
70
71                 self.anomalies_injected.append({
72                     'type': 'extreme_jump',
```

```
62         'index': i,
63         'magnitude': jump_magnitude * 100,
64         'timestamp': time.time()
65     })
66
67     yield {
68         'asset_id': 'BTC',
69         'timestamp': time.time(),
70         'price': price_to_send,
71         'volume': np.random.uniform(1000, 1000000),
72         'bid': price_to_send * 0.999 if price_to_send else
73         None,
74         'ask': price_to_send * 1.001 if price_to_send else
75         None
76     }
77
78     time.sleep(0.001) # 1ms entre puntos
79
80     def get_anomaly_report(self) -> Dict:
81         """Reporte de anomalias inyectadas"""
82         return {
83             'total_anomalies': len(self.anomalies_injected),
84             'null_values': len([a for a in self.anomalies_injected if
85                 a['type'] == 'null']),
86             'extreme_jumps': len([a for a in self.anomalies_injected if
87                 a['type'] == 'extreme_jump']),
88             'anomalies': self.anomalies_injected
89         }
90
91 class DataIngestionPipeline:
92     """Pipeline de ingesta con validacion, outlier detection y
93     limpieza"""
94
95     def __init__(self, db_connection, queue_size: int = 10000):
96         self.db = db_connection
97         self.queue = asyncio.Queue(maxsize=queue_size)
98         self.stats = {
99             'processed': 0,
100            'valid': 0,
101            'filtered_outliers': 0,
102            'filtered_nulls': 0,
103            'deduplicated': 0,
104            'stored': 0
105        }
```

```

100
101         }
102
103     self.last_seen = {} # Para deduplicacion
104
105     async def ingest_from_stream(self, stream: Generator):
106         """Consume stream y coloca en queue"""
107         for record in stream:
108             try:
109                 await asyncio.wait_for(
110                     self.queue.put(record),
111                     timeout=1.0
112                 )
113             except asyncio.TimeoutError:
114                 print(f"Queue llena, descartando record")
115
116     async def process_queue(self):
117         """Procesa items de queue con validacion y limpieza"""
118         while True:
119             try:
120                 record = await asyncio.wait_for(
121                     self.queue.get(),
122                     timeout=5.0
123                 )
124
125                 self.stats['processed'] += 1
126
127                 # Step 1: Validacion
128                 if not self._validate_record(record):
129                     self.stats['filtered_nulls'] += 1
130                     continue
131
132                 # Step 2: Deteccion de outliers
133                 if self._is_outlier(record):
134                     self.stats['filtered_outliers'] += 1
135                     continue
136
137                 # Step 3: Deduplicacion
138                 if self._is_duplicate(record):
139                     self.stats['deduplicated'] += 1
140                     continue
141
142                 # Step 4: Almacenamiento
143                 self._store_record(record)
144                 self.stats['valid'] += 1

```

```
143             self.stats['stored'] += 1
144
145         except asyncio.TimeoutError:
146             break
147
148     def _validate_record(self, record: Dict) -> bool:
149         """Valida integridad basica"""
150         if record.get('price') is None:
151             return False
152         if record.get('timestamp') is None:
153             return False
154         if not isinstance(record.get('price'), (int, float)):
155             return False
156         return True
157
158     def _is_outlier(self, record: Dict) -> bool:
159         """Detecta outliers usando metodo IQR"""
160         # Implementacion simplificada
161         price = record['price']
162
163         # Rechazar precios negativos
164         if price < 0:
165             return True
166
167         # Rechazar precios extremos
168         if price > 1e6:
169             return True
170
171         # Comparar con historico
172         asset_id = record['asset_id']
173         if asset_id in self.last_seen:
174             last_price = self.last_seen[asset_id]['price']
175             price_change = abs(price - last_price) / last_price
176
177             # Rechazar cambios >500%
178             if price_change > 5.0:
179                 return True
180
181         return False
182
183     def _is_duplicate(self, record: Dict) -> bool:
184         """Previene registros duplicados"""
185         key = (record['asset_id'], round(record['timestamp'], 1))
```

```
186
187     if key in self.last_seen:
188         # Si vimos este timestamp hace poco, es probable
189         # duplicado
190
191         self.last_seen[record['asset_id']] = record
192
193     return False
194
195 def _store_record(self, record: Dict):
196     """Almacena en TSDB (stub)"""
197     # En produccion: InfluxDB, QuestDB, etc.
198     pass
199
200 def get_stats(self) -> Dict:
201     """Retorna estadisticas de procesamiento"""
202     return {
203         **self.stats,
204         'acceptance_rate': self.stats['valid'] / max(self.stats['processed']), 1) * 100
205     }
```

Listing 5.1: Simulador de API con volatilidad realista

5.1.2. OHLC Aggregator con Tolerancia a Fallos

```
18     """
19     Calcula OHLC para un minuto
20
21     Args:
22         asset_id: ID del activo
23         minute: Timestamp de minuto (inicio)
24         prices: Series de precios del minuto
25     """
26
27     if prices.empty:
28         self.logger.warning(f"Minuto {minute} sin datos para {asset_id}")
29         return None
30
31     ohlc_data = {
32         'asset_id': asset_id,
33         'timestamp': minute,
34         'open': float(prices.iloc[0]),
35         'high': float(prices.max()),
36         'low': float(prices.min()),
37         'close': float(prices.iloc[-1]),
38         'volume': len(prices), # Count de ticks
39         'vwap': self._calculate_vwap(prices)
40     }
41
42     return ohlc_data
43
44     @staticmethod
45     def _calculate_vwap(prices: pd.Series) -> float:
46         """Volume-weighted average price"""
47         # Implementacion: necesitaria volumenes
48         return float(prices.mean())
49
50     def store_ohlc(self, ohlc_data: Dict):
51         """Almacena OHLC con fallback"""
52         try:
53             # Intentar almacenar en TSDB principal
54             self._insert_tsdb(ohlc_data)
55         except Exception as e:
56             self.logger.error(f"Error en TSDB principal: {str(e)}")
57             # Fallback: almacenar en SQLite local
58             self._insert_sqlite_fallback(ohlc_data)
```

```
60     def _insert_tsdb(self, ohlc_data: Dict):
61         """Inserta en TimeSeries DB (InfluxDB, QuestDB)"""
62         # Stub para ejemplo
63         pass
64
65     def _insert_sqlite_fallback(self, ohlc_data: Dict):
66         """Fallback a SQLite para recuperacion"""
67         cursor = self.db.cursor()
68         cursor.execute("""
69             INSERT INTO ohlc_fallback
70                 (asset_id, timestamp, open, high, low, close, volume)
71             VALUES (?, ?, ?, ?, ?, ?, ?)
72         """, (
73             ohlc_data['asset_id'],
74             ohlc_data['timestamp'].isoformat(),
75             ohlc_data['open'],
76             ohlc_data['high'],
77             ohlc_data['low'],
78             ohlc_data['close'],
79             ohlc_data['volume']
80         ))
81         self.db.commit()
```

Listing 5.2: Agregador OHLC resiliente

5.2. Pregunta 4: Pipeline Tolerante a Fallos

5.2.1. Resiliencia Mediante Pattern: Circuit Breaker

```
1 from enum import Enum
2 import time
3 from typing import Callable, Any
4
5 class CircuitState(Enum):
6     CLOSED = "closed"          # Normal operation
7     OPEN = "open"              # Failing, reject
8     HALF_OPEN = "half_open"    # Testing recovery
9
10 class CircuitBreaker:
11     """Circuit breaker con fallback automatico"""
12
13     def __init__(self,
```

```

14         failure_threshold: int = 5,
15         recovery_timeout: int = 60,
16         expected_exception: Exception = Exception):
17     """
18
19     Args:
20         failure_threshold: Numero de fallos antes de abrir
21             recovery_timeout: Segundos antes de intentar recuperacion
22             expected_exception: Tipo de excepcion a capturar
23     """
24
25     self.failure_threshold = failure_threshold
26     self.recovery_timeout = recovery_timeout
27     self.expected_exception = expected_exception
28
29     self.failure_count = 0
30     self.last_failure_time = None
31     self.state = CircuitState.CLOSED
32
33     def call(self, func: Callable, *args, fallback: Callable = None,
34             **kwargs) -> Any:
35         """
36
37         Ejecuta funcion con proteccion de circuit breaker
38
39         Args:
40             func: Funcion a ejecutar
41             fallback: Funcion alternativa si esta abierto
42         """
43
44
45         if self.state == CircuitState.OPEN:
46             if self._should_attempt_reset():
47                 self.state = CircuitState.HALF_OPEN
48                 print(f"[CircuitBreaker] Intentando recuperacion")
49             else:
50                 if fallback:
51                     print(f"[CircuitBreaker] OPEN - usando fallback")
52                     return fallback(*args, **kwargs)
53                 raise Exception("Circuit breaker is OPEN")
54
55         try:
56             result = func(*args, **kwargs)
57             self._on_success()
58             return result
59
60         except self.expected_exception as e:

```

```

56         self._on_failure()
57     if fallback:
58         print(f"[CircuitBreaker] Fallback despues de fallo: {str(e)}")
59         return fallback(*args, **kwargs)
60     raise
61
62     def _on_success(self):
63         """Registra exito y resetea counters"""
64         self.failure_count = 0
65         self.state = CircuitState.CLOSED
66
67     def _on_failure(self):
68         """Registra fallo y evalua apertura"""
69         self.failure_count += 1
70         self.last_failure_time = time.time()
71
72         if self.failure_count >= self.failure_threshold:
73             self.state = CircuitState.OPEN
74             print(f"[CircuitBreaker] OPENED despues de {self.failure_count} fallos")
75
76     def _should_attempt_reset(self) -> bool:
77         """Evalua si es hora de intentar recuperacion"""
78         if self.last_failure_time is None:
79             return False
80
81         elapsed = time.time() - self.last_failure_time
82         return elapsed >= self.recovery_timeout
83
84 class FaultTolerantPipeline:
85     """Pipeline con recuperacion automatica y fallbacks"""
86
87     def __init__(self):
88         self.primary_breaker = CircuitBreaker(
89             failure_threshold=3,
90             recovery_timeout=60
91         )
92         self.backup_breaker = CircuitBreaker(
93             failure_threshold=5,
94             recovery_timeout=120
95         )
96

```

```
97     def process_with_failover(self, data):
98         """Procesa con multiples niveles de failover"""
99
100        # Level 1: Intenta primario
101        try:
102            return self.primary_breaker.call(
103                self._primary_processor,
104                data,
105                fallback=self._backup_processor
106            )
107        except Exception as e:
108            print(f"[Pipeline] Fallo primario: {str(e)}")
109
110        # Level 2: Backup explicito
111        return self.backup_breaker.call(
112            self._backup_processor,
113            data,
114            fallback=self._local_cache_fallback
115        )
116
117    def _primary_processor(self, data):
118        """Procesador primario (e.g., TSDB remoto)"""
119        # Simulacion: puede fallar
120        if np.random.random() < 0.1: # 10% prob. fallo
121            raise ConnectionError("TSDB timeout")
122        return "processed_primary"
123
124    def _backup_processor(self, data):
125        """Procesador backup (e.g., SQLite local)"""
126        # Implementacion de fallback
127        return "processed_backup"
128
129    def _local_cacheFallback(self, data):
130        """Fallback final: cache local"""
131        return "processed_local_cache"
```

Listing 5.3: Circuit Breaker pattern para pipelines

Capítulo 6

Mecanismos de Seguridad: Preguntas Bonus

6.1. ML Engineer Bonus: Safety Mechanism contra Volatilidad Extrema

6.1.1. Framework: Adaptive Risk Management

Cuando el modelo genera predicciones en mercado 10x más volátil y pierde dinero drásticamente, se requiere:

1. **Early Warning System:** Detección de cambio de régimen de mercado
2. **Automatic Position Sizing:** Reducción dinámica de exposición
3. **Circuit Breaker:** Halt automático de trading

```
1 import numpy as np
2 from dataclasses import dataclass
3 from enum import Enum
4
5 class MarketRegime(Enum):
6     NORMAL = "normal"
7     ELEVATED = "elevated"
8     CRISIS = "crisis"
9
10 @dataclass
11 class RiskMetrics:
12     vix_equivalent: float # Volatilidad implicita
```

```
13     drawdown_current: float # Drawdown actual
14     loss_rate: float # Tasa de perdida
15     prediction_confidence: float # Confianza del modelo
16
17 class SafetyBreakerSystem:
18     """Sistema que detiene trading automaticamente ante volatilidad
19     extrema"""
20
21     def __init__(self):
22         self.regime = MarketRegime.NORMAL
23         self.regime_history = []
24         self.max_drawdown_threshold = 0.10 # 10% max drawdown
25         self.position_multiplier = 1.0
26         self.trading_enabled = True
27
28         # Thresholds de cambio de regimen
29         self.regime_thresholds = {
30             'normal_to_elevated': 0.03,           # 3x volatilidad
31             'elevated_to_crisis': 0.1,          # 10x volatilidad
32             'crisis_recovery': 0.05            # 5x para recuperacion
33         }
34
35     def evaluate_market_conditions(self,
36                                     current_metrics: RiskMetrics) ->
37     Dict:
38         """
39             Evalua condiciones de mercado y toma decisiones automaticas
40
41             Returns:
42                 {
43                     'regime': MarketRegime,
44                     'action': 'continue' | 'reduce_position' | '
45                     halt_trading',
46                     'position_multiplier': float,
47                     'reason': str
48                 }
49         """
50
51     # Step 1: Deteccion de cambio de regimen
52     new_regime = self._classify_regime(current_metrics)
```

```
51     if new_regime != self.regime:
52         self.regime = new_regime
53         self.regime_history.append({
54             'timestamp': datetime.now(),
55             'old_regime': self.regime,
56             'new_regime': new_regime,
57             'trigger_metrics': current_metrics
58         })
59         print(f"[REGIME CHANGE] {self.regime.name}")
60
61     # Step 2: Evaluacion de metricas de riesgo
62     action = 'continue'
63     reason = 'Operacion normal'
64
65     if current_metrics.drawdown_current > self.
max_drawdown_threshold:
66         self.trading_enabled = False
67         action = 'halt_trading'
68         reason = f"Drawdown {current_metrics.drawdown_current
:.1%} > threshold"
69
70     elif self.regime == MarketRegime.CRISIS:
71         self.position_multiplier = 0.2 # 20% posicion
72         action = 'reduce_position'
73         reason = "Regimen de crisis detectado"
74
75     elif self.regime == MarketRegime.ELEVATED:
76         self.position_multiplier = 0.5 # 50% posicion
77         action = 'reduce_position'
78         reason = "Volatilidad elevada"
79
80     else:
81         self.position_multiplier = 1.0
82         action = 'continue'
83
84     # Step 3: Validacion de confianza del modelo
85     if current_metrics.prediction_confidence < 0.6:
86         action = 'reduce_position'
87         reason = f"Baja confianza: {current_metrics.
prediction_confidence:.1%}"
88
89     return {
90         'regime': self.regime,
```

```

91         'action': action,
92         'position_multiplier': self.position_multiplier,
93         'reason': reason,
94         'trading_enabled': self.trading_enabled
95     }
96
97     def _classify_regime(self, metrics: RiskMetrics) -> MarketRegime:
98         """Clasifica regimen de mercado basado en volatilidad"""
99
100        # VIX-like metric: ratio de volatilidad respecto a baseline
101        volatility_ratio = metrics.vix_equivalent
102
103        if volatility_ratio > self.regime_thresholds[',
104            elevated_to_crisis']:
105            return MarketRegime.CRISIS
106
107        elif volatility_ratio > self.regime_thresholds[',
108            normal_to_elevated']:
109            return MarketRegime.ELEVATED
110
111        else:
112            return MarketRegime.NORMAL
113
114    def get_position_adjustment(self,
115                                base_position_size: float) -> float:
116        """Ajusta tamano de posicion segun condiciones"""
117
118        if not self.trading_enabled:
119            return 0.0
120
121        return base_position_size * self.position_multiplier
122
123    def get_system_status(self) -> Dict:
124        """Status completo del sistema"""
125
126        return {
127            'current_regime': self.regime.value,
128            'position_multiplier': self.position_multiplier,
129            'trading_enabled': self.trading_enabled,
130            'regime_history_length': len(self.regime_history),
131            'recent_transitions': self.regime_history[-5:] if self.
132            regime_history else []
133        }

```

Listing 6.1: Safety system contra volatilidad extrema

6.2. Data Engineer Bonus: Consenso Multi-proveedor

6.2.1. Algoritmo de Reconciliación

Cuando recibimos datos de 3 proveedores con discrepancias:

$$P_{\text{consenso}} = \operatorname{argmin}_p \sum_i w_i |p - p_i| \quad (6.1)$$

Donde w_i es peso del proveedor i (basado en histórico de confiabilidad).

```

1 import numpy as np
2 from typing import Dict, List
3 from scipy.stats import median_abs_deviation as mad
4
5 class MultiProviderPriceConsensus:
6     """Reconcilia precios de multiples proveedores"""
7
8     def __init__(self, providers: List[str]):
9         """
10            Args:
11                providers: Lista de IDs de proveedores
12            """
13
14        self.providers = providers
15        self.provider_stats = {
16            p: {'errors': 0, 'total': 0, 'reliability': 1.0}
17            for p in providers
18        }
19        self.consensus_history = []
20
21    def reconcile_price(self,
22                        prices: Dict[str, float],
23                        timestamp: str) -> Dict:
24        """
25            Reconcilia precios de multiples proveedores
26
27            Args:
28                prices: {'provider1': 100.5, 'provider2': 100.45, ,
29                'provider3': 100.6}
30
31            Returns:
32                {
33                    'consensus_price': 100.5,
34                    'method': 'weighted_median',
35                    'confidence': 0.95,

```

```
34         'outliers': ['provider3'],
35         'reasoning': str
36     }
37 """
38
39 # Step 1: Validacion inicial
40 valid_prices = {p: v for p, v in prices.items() if v > 0}
41
42 if len(valid_prices) == 0:
43     return self._handle_no_valid_prices()
44
45 if len(valid_prices) == 1:
46     provider, price = list(valid_prices.items())[0]
47     return {
48         'consensus_price': price,
49         'method': 'single_provider',
50         'confidence': 0.5, # Baja confianza con un proveedor
51         'outliers': [],
52         'reasoning': f'Solo {provider} tiene precio valido'
53     }
54
55 # Step 2: Deteccion de outliers con MAD (Median Absolute
56 Deviation)
57 price_values = np.array(list(valid_prices.values()))
58 median_price = np.median(price_values)
59 deviations = np.abs(price_values - median_price)
60 mad_value = mad(price_values)
61
62 # Threshold: |precio - mediana| > 3*MAD
63 outlier_threshold = 3 * mad_value
64
65 outliers = []
66 inliers = {}
67
68 for provider, price in valid_prices.items():
69     deviation = abs(price - median_price)
70
71     if deviation > outlier_threshold and len(valid_prices) >
72     2:
73         outliers.append(provider)
74     else:
75         inliers[provider] = price
```

```
75     # Step 3: Calculo de consenso
76     if len(inliers) == 0:
77         # Todos son outliers, usar mediana completa
78         consensus_price = median_price
79         method = 'median_all'
80     else:
81         # Usar weighted median de inliers
82         inlier_prices = np.array(list(inliers.values()))
83         inlier_providers = list(inliers.keys())
84
85         # Pesos basados en confiabilidad historica
86         weights = np.array([
87             self.provider_stats[p]['reliability']
88             for p in inlier_providers
89         ])
90
91         weights = weights / weights.sum()
92
93         # Weighted median
94         consensus_price = self._weighted_median(inlier_prices,
95         weights)
96         method = 'weighted_median_inliers'
97
98         # Step 4: Calculo de confianza
99         confidence = self._calculate_confidence(
100             valid_prices, consensus_price, outliers
101         )
102
103         # Step 5: Actualizar estadisticas de proveedores
104         self._update_provider_stats(valid_prices, consensus_price,
105         outliers)
106
107         result = {
108             'consensus_price': float(consensus_price),
109             'method': method,
110             'confidence': float(confidence),
111             'outliers': outliers,
112             'outlier_reasons': [
113                 f'{p}: {abs(valid_prices[p] - median_price)/
114 median_price*100:.2f} %'
115                 f"del mediana"
116                 for p in outliers
117             ],
118         },
```

```
115     'provider_stats': {
116         p: self.provider_stats[p]
117         for p in self.providers
118     }
119 }
120
121     self.consensus_history.append({
122         'timestamp': timestamp,
123         'result': result
124     })
125
126     return result
127
128 @staticmethod
129 def _weighted_median(values: np.ndarray, weights: np.ndarray) -> float:
130     """Calcula mediana ponderada"""
131     sorted_indices = np.argsort(values)
132     sorted_values = values[sorted_indices]
133     sorted_weights = weights[sorted_indices]
134
135     cumsum_weights = np.cumsum(sorted_weights)
136     target_weight = 0.5
137
138     idx = np.argmax(cumsum_weights >= target_weight)
139     return sorted_values[idx]
140
141 def _calculate_confidence(self,
142                           prices: Dict[str, float],
143                           consensus: float,
144                           outliers: List[str]) -> float:
145     """Calcula confianza en consenso"""
146
147     # Menos outliers = mas confianza
148     outlier_ratio = len(outliers) / len(prices)
149
150     # Menos dispersion = mas confianza
151     price_values = np.array(list(prices.values()))
152     coefficient_variation = np.std(price_values) / np.mean(
153     price_values)
154
155     # Combinar factores
156     confidence = (1 - outlier_ratio) * (1 - min(
```

```
    coefficient_variation, 0.1))

156
157     return max(0.1, min(1.0, confidence))

158
159     def _update_provider_stats(self,
160                               prices: Dict[str, float],
161                               consensus: float,
162                               outliers: List[str]):
163         """Actualiza confiabilidad de proveedores"""
164
165         for provider, price in prices.items():
166             self.provider_stats[provider]['total'] += 1
167
168             if provider in outliers:
169                 self.provider_stats[provider]['errors'] += 1
170
171             # Confiabilidad = (1 - error_rate)
172             error_rate = self.provider_stats[provider]['errors'] / \
173                         max(self.provider_stats[provider]['total'],
174                             1)
175
176             self.provider_stats[provider]['reliability'] = 1 - min(
177                 error_rate, 1.0)
178
179     def _handle_no_valid_prices(self) -> Dict:
180         """Maneja caso sin precios validos"""
181         return {
182             'consensus_price': None,
183             'method': 'no_valid_data',
184             'confidence': 0.0,
185             'outliers': list(self.providers),
186             'reasoning': 'Todos los proveedores retornaron precios
187 invalidos'
188         }
```

Listing 6.2: Sistema de consenso multi-proveedor

Capítulo 7

Conclusión: Arquitectura Integrada

7.1. Sistema Completo: Flujo End-to-End

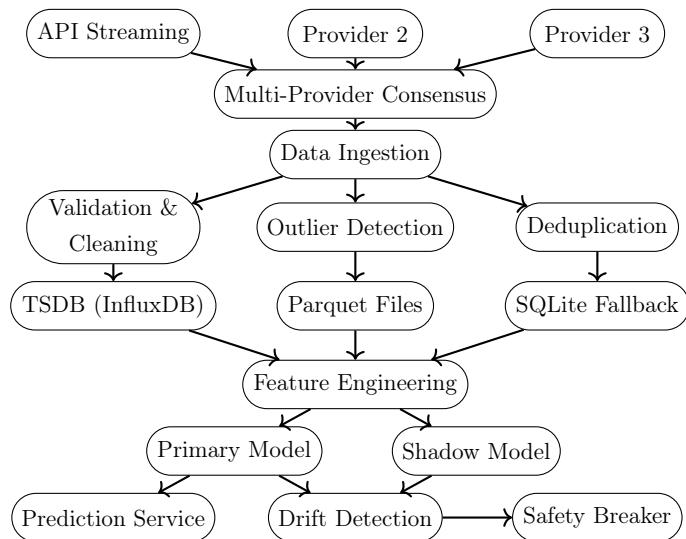


Figura 7.1: Arquitectura integrada del sistema de predicción (Corregida)

7.2. Key Insights y Recomendaciones

7.2.1. Diseño para Producción

1. **Separación de concernos:** Data layer, ML layer, API layer están desacoplados
2. **Resiliencia multi-nivel:** Circuit breakers, fallbacks, caches locales
3. **Observabilidad:** Logging exhaustivo, métricas, alertas automáticas
4. **Validación en todos los niveles:** Pydantic, schema validation, sanity checks

7.2.2. Métricas Críticas de Monitoreo

Métrica	Target	Acción
Latencia P99	< 150ms	Alertar si > 180ms
SLA Violation	< 5 %	Invocar escalación
Model Drift Score	< 0.15	Retrain si > 0.2
Data Gap Duration	< 10 min	Investigar si > 30min
Deduplication Rate	< 1 %	Review si > 2 %
Provider Consensus	> 0.9	Flag si < 0.8

Cuadro 7.1: SLAs y umbrales de monitoreo

7.3. Roadmap de Implementación

1. **Fase 1 (Semanas 1-2):** Implementar validadores Pydantic y tests unitarios
2. **Fase 2 (Semanas 3-4):** Containerización Docker y CI/CD pipeline
3. **Fase 3 (Semanas 5-6):** Data pipeline con TSDB y fallbacks
4. **Fase 4 (Semanas 7-8):** Drift detection y retraining automático
5. **Fase 5 (Semanas 9-10):** Safety systems y circuit breakers
6. **Fase 6 (Semanas 11-12):** Load testing, optimization, deployment

Apéndice A

Referencias Técnicas

A.1. Librerías Recomendadas

- **ML:** scikit-learn, XGBoost, ONNX, SHAP
- **Data:** pandas, numpy, pyarrow, DuckDB
- **API:** FastAPI, Pydantic, uvicorn
- **Storage:** InfluxDB, QuestDB, SQLite
- **Monitoring:** Prometheus, Grafana, ELK Stack
- **Testing:** pytest, hypothesis, locust

A.2. Estándares de Código

- Linting: pylint, flake8, black
- Type hints: mypy
- Documentation: Sphinx, docstrings
- Versionamiento: Git flow